

Step 1: Import Required Libraries

```
In [1]: # Import necessary libraries
```

```
import pickle
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from scipy.stats import randint
from datetime import timedelta
from sklearn.model_selection import train_test_split
from sklearn.feature_selection import SelectKBest, f_regression
from sklearn.preprocessing import StandardScaler
from xgboost import XGBRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

Step 2: Load Data

```
In [2]: # Load climate data
```

```
folder_path = r'D:/Pasture Growth/datasets/'

climate_data = pd.concat([
    pd.read_csv(f'{folder_path}climate_{year}.csv') for year in range(2017,
])

climate_data.head()
```

Out[2]:

		PADDock_ID	DATE	RAIN	MAX_TEMP	MIN
0	E5BEE111190304432B4EEFB0B7FA23A7C23CF8960E4C7F...		2017-03-22	50.2	28.4	
1	4C2EB78791FD1F9A42113DA94A14C1190BD8FBB8E87F96...		2017-03-22	52.5	28.5	
2	9C2AE9804DFBFDF6791C03D9FCE5DCBF9DA6AE7476477...		2017-03-22	52.5	28.5	
3	04C1E345BBF76EA3E2B076C7E1E17290AD77F6BD159AD7...		2017-03-22	52.5	28.5	
4	9F7943EC335457BE4BF1A456EC8AE9567902A02B4173C8...		2017-03-22	52.5	28.5	

```
In [3]: climate_data.shape
```

```
Out[3]: (33276908, 9)
```

```
In [4]: # Load TSDM data  
  
tsdm_data = pd.read_csv(r'D:\Pasture Growth\datasets\nrm1010_tsdm.csv')  
  
tsdm_data.head()
```

```
Out[4]:
```

	PADDOCK_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	5/01/2017	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	20/01/2017	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	4/02/2017	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	19/02/2017	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	6/03/2017	2330.000000

```
In [5]: tsdm_data.shape
```

```
Out[5]: (1030320, 3)
```

```
In [6]: # Load paddock metadata  
  
paddock_metadata = pd.read_csv(r'D:\Pasture Growth\datasets\paddock_metadata.csv')  
  
paddock_metadata.head()
```

```
Out[6]:
```

	PADDOCK_ID	CROP_TYPE	LANDSIZE_HA	PA
0	0AE10585CD37DE9A5DF983CF95D7237D534E6CE7EB6081...	natural grasses	29.90	
1	ED2AAAB2838D848AC4E7BDE8BA13941B7105ECF65DF3F4...	natural grasses	16.15	
2	435141E0BE1DBFD4E37DC6E374B4BD4C6D91CFC722E73A...	natural grasses	8.27	
3	AD6846319B84A902CB0B71B3F640643FCDA283D84B2DE8...	natural grasses	22.07	
4	7A85B8453613D37697502D75F3DBD922FBEE91FC6AB8A8...	natural grasses	0.82	

```
In [7]: paddock_metadata.shape
```

```
Out[7]: (17966, 5)
```

Step 3: Data Preprocessing & EDA

```
In [8]: # Convert DATE to datetime format
```

```
climate_data['DATE'] = pd.to_datetime(climate_data['DATE'], format='%Y-%m-%d')

tsdm_data['OBSERVATION_DATE'] = pd.to_datetime(tsdm_data['OBSERVATION_DATE'])
```

```
In [9]: # Check unique values in PASTURE_STATE column

print(paddock_metadata['PASTURE_STATE'].unique())

['Grazing' 'Laneway' 'Yard' 'Cropping' 'Feedlot' 'Vegetation' 'Hay' 'Pen']
```

```
In [10]: # Filter the DataFrame where PASTURE_STATE is 'Grazing'

paddock_metadata = paddock_metadata[paddock_metadata['PASTURE_STATE'] == 'Gr']
```

```
In [11]: print(paddock_metadata['PASTURE_STATE'].unique())

['Grazing']
```

```
In [12]: paddock_metadata.shape
```

```
Out[12]: (14956, 5)
```

```
In [13]: # Convert PADDOCK_ID to lowercase for consistency

tsdm_data['PADDOCK_ID'] = tsdm_data['PADDOCK_ID'].str.lower()
climate_data['PADDOCK_ID'] = climate_data['PADDOCK_ID'].str.lower()
paddock_metadata['PADDOCK_ID'] = paddock_metadata['PADDOCK_ID'].str.lower()

# Strip spaces from PADDOCK_ID

tsdm_data['PADDOCK_ID'] = tsdm_data['PADDOCK_ID'].str.strip()
climate_data['PADDOCK_ID'] = climate_data['PADDOCK_ID'].str.strip()
paddock_metadata['PADDOCK_ID'] = paddock_metadata['PADDOCK_ID'].str.strip()

# Common PADDOCK_ID values across all three datasets

common_ids_tsdm_climate = set(tsdm_data['PADDOCK_ID']).intersection(set(clim
common_ids_all = common_ids_tsdm_climate.intersection(set(paddock_metadata['

# Number of common PADDOCK_IDs

print(f"Number of common PADDOCK_IDs between TSDM and Climate data: {len(com
print(f"Number of common PADDOCK_IDs across all datasets: {len(common_ids_all
```

```
Number of common PADDOCK_IDs between TSDM and Climate data: 4855
Number of common PADDOCK_IDs across all datasets: 4808
```

```
In [14]: # Merge TSDM data with Climate data

merged_data = pd.merge(tsdm_data, climate_data,
                      left_on=['PADDOCK_ID', 'OBSERVATION_DATE'],
                      right_on=['PADDOCK_ID', 'DATE'],
                      how='inner')

# Merge with paddock metadata
```

```
merged_data = pd.merge(merged_data, paddock_metadata,
                      on='PADDOCK_ID',
                      how='inner')

merged_data.head()
```

Out[14]:

	PADDOCK_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

In [15]: merged_data.shape

Out[15]: (701968, 15)

In [16]: *# Remove the DATE and CREATION_DATE column from the merged_data*

```
merged_data.drop(columns=['DATE', 'CREATION_DATE'], inplace=True)

merged_data.columns
```

Out[16]: Index(['PADDOCK_ID', 'OBSERVATION_DATE', 'TSDM_MEAN', 'RAIN', 'MAX_TEMP',
 'MIN_TEMP', 'RH_TMAX', 'RH_TMIN', 'EVAP', 'RADIATION', 'CROP_TYPE',
 'LANDSIZE_HA', 'PASTURE_STATE'],
 dtype='object')

In [17]: merged_data.head()

	PADDOCK_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

In [18]: merged_data.shape

Out[18]: (701968, 13)

In [19]: merged_data.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 701968 entries, 0 to 701967
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   PADDOCK_ID       701968 non-null   object  
 1   OBSERVATION_DATE 701968 non-null   datetime64[ns] 
 2   TSMD_MEAN        687426 non-null   float64 
 3   RAIN              701968 non-null   float64 
 4   MAX_TEMP          701968 non-null   float64 
 5   MIN_TEMP          701968 non-null   float64 
 6   RH_TMAX           701968 non-null   float64 
 7   RH_TMIN           701968 non-null   float64 
 8   EVAP              701968 non-null   float64 
 9   RADIATION         701968 non-null   float64 
 10  CROP_TYPE         701092 non-null   object  
 11  LANDSIZE_HA       701968 non-null   float64 
 12  PASTURE_STATE     701968 non-null   object  
dtypes: datetime64[ns](1), float64(9), object(3)
memory usage: 69.6+ MB
```

```
In [20]: # Count of missing values in each column
merged_data.isnull().sum()
```

```
Out[20]: PADDOCK_ID      0
OBSERVATION_DATE    0
TSMD_MEAN          14542
RAIN                0
MAX_TEMP            0
MIN_TEMP            0
RH_TMAX             0
RH_TMIN             0
EVAP                0
RADIATION           0
CROP_TYPE           876
LANDSIZE_HA         0
PASTURE_STATE       0
dtype: int64
```

```
In [21]: # Check for outlier in TSMD_MEAN column
plt.figure(figsize=(10, 6))
sns.boxplot(x=merged_data['TSMD_MEAN'])
plt.title('Boxplot of TSMD_MEAN')
plt.show()
```

Boxplot of TSMD_MEAN



```
In [22]: # Impute missing TSMD_MEAN values with the median  
merged_data['TSMD_MEAN'].fillna(merged_data['TSMD_MEAN'].median(), inplace=True)
```

```
In [23]: # Count of missing values in each column  
merged_data.isnull().sum()
```

```
Out[23]: PADDock_ID      0  
OBSERVATION_DATE    0  
TSMD_MEAN          0  
RAIN                0  
MAX_TEMP           0  
MIN_TEMP           0  
RH_TMAX             0  
RH_TMIN             0  
EVAP                0  
RADIATION          0  
CROP_TYPE          876  
LANDSIZE_HA         0  
PASTURE_STATE       0  
dtype: int64
```

```
In [24]: # Impute missing CROP_TYPE values with the mode  
merged_data['CROP_TYPE'].fillna(merged_data['CROP_TYPE'].mode()[0], inplace=True)
```

```
In [25]: # Count of missing values in each column
```

```
merged_data.isnull().sum()
```

```
Out[25]: PADDock_ID      0  
OBSERVATION_DATE    0  
TSDM_MEAN          0  
RAIN                0  
MAX_TEMP           0  
MIN_TEMP           0  
RH_TMAX             0  
RH_TMIN             0  
EVAP                0  
RADIATION          0  
CROP_TYPE          0  
LANDSIZE_HA         0  
PASTURE_STATE       0  
dtype: int64
```

```
In [26]: # check for stats
```


```
merged_data.describe()
```

```
Out[26]:   OBSERVATION_DATE    TSDM_MEAN      RAIN      MAX_TEMP      MIN_TEMP  
count            701968  701968.000000  701968.000000  701968.000000  701968.000000  
mean  2019-12-28 12:00:00    1720.159833    1.979314    21.352619    8.203790  
min   2017-01-05 00:00:00    41.833333    0.000000    2.500000   -6.000000  
25%   2018-06-29 00:00:00   1151.583333    0.000000   15.500000    3.400000  
50%   2019-12-28 12:00:00   1708.714286    0.000000   20.700000    7.400000  
75%   2021-06-28 00:00:00   2223.285714    0.800000   26.900000   13.400000  
max   2022-12-20 00:00:00   7532.500000   70.800000   45.600000   27.000000  
std        NaN      695.954315     5.666996     7.597632     6.186689
```

```
In [27]: # Box Plot of TSDM_MEAN
```

```
plt.figure(figsize=(10, 5))  
sns.boxplot(x=merged_data['TSDM_MEAN'])  
plt.title('Box Plot of TSDM_MEAN')  
plt.xlabel('TSDM_MEAN')  
plt.show()
```


Box Plot of TSDM_MEAN



In [28]: # Box Plot of RAIN

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['RAIN'])
plt.title('Box Plot of RAIN')
plt.xlabel('RAIN')
plt.show()
```

Box Plot of RAIN




In [29]: # Box Plot of MAX_TEMP

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['MAX_TEMP'])
```

```
plt.title('Box Plot of MAX_TEMP')
plt.xlabel('MAX_TEMP')
plt.show()
```


Box Plot of MAX_TEMP



In [30]: # Box Plot of MIN_TEMP


```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['MIN_TEMP'])
plt.title('Box Plot of MIN_TEMP')
plt.xlabel('MIN_TEMP')
plt.show()
```

Box Plot of MIN_TEMP



```
In [31]: # Box Plot of RH_TMAX
```


```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['RH_TMAX'])
plt.title('Box Plot of RH_TMAX')
plt.xlabel('RH_TMAX')
plt.show()
```



```
In [32]: # Box Plot of RH_TMIN
```

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['RH_TMIN'])
plt.title('Box Plot of RH_TMIN')
plt.xlabel('RH_TMIN')
plt.show()
```


Box Plot of RH_TMIN



```
In [33]: # Box Plot of EVAP
```

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['EVAP'])
plt.title('Box Plot of EVAP')
plt.xlabel('EVAP')
plt.show()
```

Box Plot of EVAP




```
In [34]: # Box Plot of RADIATION
```

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['RADIATION'])
```

```
plt.title('Box Plot of RADIATION')
plt.xlabel('RADIATION')
plt.show()
```


Box Plot of RADIATION



In [35]: # Box Plot of LANDSIZE_HA

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=merged_data['LANDSIZE_HA'])
plt.title('Box Plot of LANDSIZE_HA')
plt.xlabel('LANDSIZE_HA')
plt.show()
```

Box Plot of LANDSIZE_HA



```
In [36]: merged_data.describe()
```

```
Out[36]:
```

	OBSERVATION_DATE	TSMD_MEAN	RAIN	MAX_TEMP	MIN_TEMP
count	701968	701968.000000	701968.000000	701968.000000	701968.000000
mean	2019-12-28 12:00:00	1720.159833	1.979314	21.352619	8.203790
min	2017-01-05 00:00:00	41.833333	0.000000	2.500000	-6.000000
25%	2018-06-29 00:00:00	1151.583333	0.000000	15.500000	3.400000
50%	2019-12-28 12:00:00	1708.714286	0.000000	20.700000	7.400000
75%	2021-06-28 00:00:00	2223.285714	0.800000	26.900000	13.400000
max	2022-12-20 00:00:00	7532.500000	70.800000	45.600000	27.000000
std	NaN	695.954315	5.666996	7.597632	6.186689

```
In [37]: merged_data.head()
```

```
Out[37]:
```

	PADDOCK_ID	OBSERVATION_DATE	TSMD_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

```
In [38]: # Function to cap outliers using IQR
```

```
def cap_outliers_iqr(merged_data):  
  
    exclude_cols = ['MIN_TEMP', 'RH_TMAX', 'RADIATION']  
  
    # Iterate over each numerical column in the DataFrame  
    for col in merged_data.select_dtypes(include=['float64']).columns:  
  
        if col in exclude_cols:  
            continue  
  
        Q1 = merged_data[col].quantile(0.25)  
        Q3 = merged_data[col].quantile(0.75)  
        IQR = Q3 - Q1  
  
        lower_bound = Q1 - 1.5 * IQR  
        upper_bound = Q3 + 1.5 * IQR  
  
        print(f"Column: {col}")  
        print(f"  Q1: {Q1}")  
        print(f"  Q3: {Q3}")  
        print(f"  IQR: {IQR}")  
        print(f"  Lower Bound: {lower_bound}")
```

```
        print(f"  Upper Bound: {upper_bound}")

    # Cap the outliers
    merged_data[col] = merged_data[col].clip(lower=lower_bound, upper=upper_bound)

    return merged_data

final_merge = cap_outliers_iqr(merged_data)
```

Column: TSDM_MEAN
Q1: 1151.583333333333
Q3: 2223.285714285714
IQR: 1071.702380952381
Lower Bound: -455.9702380952383
Upper Bound: 3830.839285714286

Column: RAIN
Q1: 0.0
Q3: 0.8
IQR: 0.8
Lower Bound: -1.2000000000000002
Upper Bound: 2.0

Column: MAX_TEMP
Q1: 15.5
Q3: 26.9
IQR: 11.39999999999999
Lower Bound: -1.599999999999979
Upper Bound: 44.0

Column: RH_TMIN
Q1: 93.3
Q3: 100.0
IQR: 6.700000000000003
Lower Bound: 83.25
Upper Bound: 110.0500000000001

Column: EVAP
Q1: 1.9
Q3: 5.3
IQR: 3.4
Lower Bound: -3.199999999999997
Upper Bound: 10.39999999999999

Column: LANDSIZE_HA
Q1: 25.749104597749998
Q3: 49.98700938925
IQR: 24.2379047915
Lower Bound: -10.607752589500002
Upper Bound: 86.3438665765

In [39]: final_merge.head()

Out[39]:

	PADDock_ID	OBSERVATION_DATE	TSMD_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

In [40]: final_merge.describe()

Out[40]:

	OBSERVATION_DATE	TSMD_MEAN	RAIN	MAX_TEMP	MIN_TEMP
count	701968	701968.000000	701968.000000	701968.000000	701968.000000
mean	2019-12-28 12:00:00	1718.284044	0.490724	21.352095	8.203790
min	2017-01-05 00:00:00	41.833333	0.000000	2.500000	-6.000000
25%	2018-06-29 00:00:00	1151.583333	0.000000	15.500000	3.400000
50%	2019-12-28 12:00:00	1708.714286	0.000000	20.700000	7.400000
75%	2021-06-28 00:00:00	2223.285714	0.800000	26.900000	13.400000
max	2022-12-20 00:00:00	3830.839286	2.000000	44.000000	27.000000
std	NaN	688.505806	0.795765	7.596043	6.186689

In [41]: # Columns to plot

```

columns_to_plot = ['TSMD_MEAN', 'RAIN', 'MAX_TEMP', 'RH_TMIN', 'EVAP', 'LAND

# Box plots and distribution plots after capping


for column in columns_to_plot:
    plt.figure(figsize=(12, 6))


    # Box Plot
    plt.subplot(1, 2, 1)
    sns.boxplot(x=final_merge[column])
    plt.title(f'Box Plot of {column} after capping')
    plt.xlabel(column)

    # Distribution Plot
    plt.subplot(1, 2, 2)
    sns.histplot(final_merge[column], kde=True, bins=30)
    plt.title(f'Distribution Plot of {column} after capping')
    plt.xlabel(column)
    plt.ylabel('Frequency')

    plt.tight_layout()
    plt.show()


```





```
In [42]: # Pairplot
```

```
selected_columns = ['TSDM_MEAN', 'RAIN', 'MAX_TEMP', 'EVAP', 'RADIATION']
sns.pairplot(final_merge[selected_columns], diag_kind='kde')
plt.suptitle('Pairplot of Climate and TSDM Variables', y=1.02)
plt.show()
```



Step 4: Feature Engineering

4.1 Date Features

```
In [43]: # Convert 'OBSERVATION_DATE' to datetime
```

```
final_merge['OBSERVATION_DATE'] = pd.to_datetime(final_merge['OBSERVATION_DA
```

```
# Extracting features from the date  
  
final_merge['Year'] = final_merge['OBSERVATION_DATE'].dt.year  
final_merge['Month'] = final_merge['OBSERVATION_DATE'].dt.month
```


In [44]: `final_merge.head()`

Out[44]:

	PADDock_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

In [45]: `# Group by year and calculate mean of TSDM_MEAN`

```
yearly_tsdm_mean = final_merge.groupby('Year')[['TSDM_MEAN']].mean()  
  
# Line plot for TSDM Mean over time yearly  
  
plt.figure(figsize=(10, 6))  
plt.plot(yearly_tsdm_mean.index, yearly_tsdm_mean.values, marker='o', color=  
plt.title('TSDM Mean Over Time (Yearly)', fontsize=14)  
plt.xlabel('Year', fontsize=12)  
plt.ylabel('Mean TSDM', fontsize=12)  
plt.grid(True)  
plt.show()
```



4.2 Ratio Features

```
In [46]: # Creating difference features
final_merge['Temp_Range'] = final_merge['MAX_TEMP'] - final_merge['MIN_TEMP']
```

```
In [47]: final_merge.head()
```

Out[47]:

	PADDOCK_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

4.3 Interaction Features

```
In [48]: # Interaction feature
final_merge['Temp_Rain_Interaction'] = final_merge['MAX_TEMP'] * final_merge
```

```
In [49]: final_merge.head()
```

Out[49]:

	PADDock_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	1585.000000
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	1620.333333
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	1850.500000
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	2303.000000
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	2330.000000

4.4 Scaling Numerical Features

In [50]:

```
# Initialize the StandardScaler  
  
scaler = StandardScaler()  
  
# List of numerical columns to scale  
  
columns_to_scale = ['TSDM_MEAN', 'RAIN', 'MAX_TEMP', 'MIN_TEMP', 'RH_TMAX',  
                    'RH_TMIN', 'EVAP', 'RADIATION', 'LANDSIZE_HA',  
                    'Temp_Range', 'Temp_Rain_Interaction']  
  
# Scale the numerical columns  
  
final_merge[columns_to_scale] = scaler.fit_transform(final_merge[columns_to_]  
print(final_merge[columns_to_scale].head())
```

	TSDM_MEAN	RAIN	MAX_TEMP	MIN_TEMP	RH_TMAX	RH_TMIN	EVAP	\
0	-0.193585	-0.616670	1.388606	1.486452	-0.772311	-1.798212	1.317150	
1	-0.142266	-0.491005	1.270123	1.615762	0.267118	0.632316	2.100872	
2	0.192033	-0.616670	2.362798	2.213821	-0.824609	-0.281171	0.822168	
3	0.849254	-0.616670	1.019993	1.082358	-1.628696	-2.099989	2.265866	
4	0.888469	-0.365339	1.112146	1.033867	-0.177418	0.632316	0.368434	

	RADIATION	LANDSIZE_HA	Temp_Range	Temp_Rain_Interaction	
0	1.473871	-0.285487	0.264898	-0.570730	
1	-0.484690	-0.285487	-0.068259	-0.378665	
2	0.924346	-0.285487	0.833226	-0.570730	
3	0.698900	-0.285487	0.206106	-0.570730	
4	0.698900	-0.285487	0.402081	-0.201470	

In [51]:

```
final_merge.head()
```

Out[51]:

	PADDock_ID	OBSERVATION_DATE	TSDM_MEAN
0	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-05	-0.193585
1	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-01-20	-0.142266
2	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-04	0.192033
3	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-02-19	0.849254
4	deaef36d983eab83c6d11ad98ffdf0daab623cc0221043...	2017-03-06	0.888469

In [52]: `final_merge.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 701968 entries, 0 to 701967
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   PADDock_ID        701968 non-null   object 
 1   OBSERVATION_DATE  701968 non-null   datetime64[ns]
 2   TSDM_MEAN         701968 non-null   float64
 3   RAIN              701968 non-null   float64
 4   MAX_TEMP          701968 non-null   float64
 5   MIN_TEMP          701968 non-null   float64
 6   RH_TMAX            701968 non-null   float64
 7   RH_TMIN            701968 non-null   float64
 8   EVAP              701968 non-null   float64
 9   RADIATION          701968 non-null   float64
 10  CROP_TYPE          701968 non-null   object 
 11  LANDSIZE_HA        701968 non-null   float64
 12  PASTURE_STATE      701968 non-null   object 
 13  Year               701968 non-null   int32  
 14  Month              701968 non-null   int32  
 15  Temp_Range          701968 non-null   float64
 16  Temp_Rain_Interaction 701968 non-null   float64
dtypes: datetime64[ns](1), float64(11), int32(2), object(3)
memory usage: 85.7+ MB
```

In [53]: `final_merge.isnull().sum()`

```
Out[53]: PADDock_ID          0  
OBSERVATION_DATE        0  
TSDM_MEAN                0  
RAIN                      0  
MAX_TEMP                 0  
MIN_TEMP                 0  
RH_TMAX                  0  
RH_TMIN                  0  
EVAP                     0  
RADIATION                0  
CROP_TYPE                0  
LANDSIZE_HA               0  
PASTURE_STATE             0  
Year                      0  
Month                     0  
Temp_Range                0  
Temp_Rain_Interaction      0  
dtype: int64
```


Step 5: Feature Selection

5.1 Correlation Matrix

```
In [54]: # Select only numeric columns for correlation analysis  
  
numeric_data = final_merge.select_dtypes(include=['int32', 'float64'])  
  
numeric_data.head()
```

```
Out[54]:   TSDM_MEAN      RAIN    MAX_TEMP    MIN_TEMP    RH_TMAX    RH_TMIN     EVAP     RADIA  
0   -0.193585 -0.616670    1.388606    1.486452   -0.772311   -1.798212   1.317150   1.47  
1   -0.142266 -0.491005    1.270123    1.615762    0.267118    0.632316   2.100872  -0.48  
2    0.192033 -0.616670    2.362798    2.213821   -0.824609   -0.281171   0.822168   0.92  
3    0.849254 -0.616670    1.019993    1.082358   -1.628696   -2.099989   2.265866   0.69  
4    0.888469 -0.365339    1.112146    1.033867   -0.177418    0.632316   0.368434   0.69
```

```
In [55]: # Calculate the correlation matrix  
  
correlation_matrix = numeric_data.corr()  
  
# Plotting the heatmap  
  
plt.figure(figsize=(12, 8))  
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", line  
plt.title('Correlation Matrix')  
plt.show()
```



```
In [56]: target_correlations = correlation_matrix['TSDM_MEAN'].sort_values(ascending=False)

print("Feature Correlations with TSDM_MEAN:")

print(target_correlations)
```

Feature Correlations with TSDM_MEAN:

TSDM_MEAN	1.000000
Year	0.519108
RH_TMIN	0.268196
RH_TMAX	0.260854
Month	0.081826
RAIN	0.045899
Temp_Rain_Interaction	0.040114
LANDSIZE_HA	-0.072778
MIN_TEMP	-0.081271
Temp_Range	-0.147897
MAX_TEMP	-0.165543
RADIATION	-0.197527
EVAP	-0.261396

Name: TSDM_MEAN, dtype: float64

Step 6: Modelling

```
In [57]: # Target variable

y = numeric_data['TSDM_MEAN'] # Target variable

# Selected features

selected_feature = [
    'Year',
    'EVAP',
    'RADIATION',
    'Temp_Range',
    'RH_TMAX',
    'MIN_TEMP',
    'MAX_TEMP',
]

X_selected_df = numeric_data[selected_feature]

# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X_selected_df, y, test_s
```

6.1 Random Forest Regressor

```
In [58]: # Initialize and train the Random Forest Regressor

rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Make predictions

y_pred = rf_model.predict(X_test)
```



```
In [59]: # Evaluate the model

mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False) # RMSE is the square root of MSE
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print evaluation metrics
print(f'Mean Squared Error (MSE): {mse}')
print(f'Root Mean Squared Error (RMSE): {rmse}')
print(f'Mean Absolute Error (MAE): {mae}')
print(f'R-squared (R2): {r2}')

Mean Squared Error (MSE): 0.2814475768773077
Root Mean Squared Error (RMSE): 0.5305163304529915
Mean Absolute Error (MAE): 0.38660633427178753
R-squared (R2): 0.7171190995832423
```

```
In [60]: # Feature importances from the trained model

importances = rf_model.feature_importances_
```

```

feature_importance_df = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': importances
})

# Sort the DataFrame by importance


feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10, 6))

# Bar plot

sns.barplot(x='Importance', y='Feature', data=feature_importance_df)
plt.title('Feature Importance from Random Forest Model')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.show()

```



6.2 XGBoost Regressor

```

In [61]: # Initialize and train the XGBoost Regressor

xgb_model = XGBRegressor(n_estimators=100, random_state=42)
xgb_model.fit(X_train, y_train)

# Make predictions

y_pred = xgb_model.predict(X_test)

```

```

In [62]: # Evaluate the model

```

```

mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False) # RMSE is the square root of MSE
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print evaluation metrics

print(f'Mean Squared Error (MSE): {mse}')
print(f'Root Mean Squared Error (RMSE): {rmse}')
print(f'Mean Absolute Error (MAE): {mae}')
print(f'R-squared (R2): {r2}')

```

Mean Squared Error (MSE): 0.33476189530268
Root Mean Squared Error (RMSE): 0.5785861174472474
Mean Absolute Error (MAE): 0.4348989154947407
R-squared (R²): 0.6635332681875447

6.3 Save Random Forest Model

In [63]: *# Save the model*

```

with open(r'D:\Pasture Growth\datasets\random_forest_model.pkl', 'wb') as file:
    pickle.dump(rf_model, file)

```

In [64]: *# Load the model*

```

with open(r'D:\Pasture Growth\datasets\random_forest_model.pkl', 'rb') as file:
    loaded_rf_model = pickle.load(file)

# Loaded model to make predictions

y_pred_loaded = loaded_rf_model.predict(X_test)

```

In [65]: *# Evaluate the loaded model*

```

mse_loaded = mean_squared_error(y_test, y_pred_loaded)
rmse_loaded = mean_squared_error(y_test, y_pred_loaded, squared=False) # RMSE is the square root of MSE
mae_loaded = mean_absolute_error(y_test, y_pred_loaded)
r2_loaded = r2_score(y_test, y_pred_loaded)

# Print evaluation metrics

print(f'Mean Squared Error (MSE): {mse_loaded}')
print(f'Root Mean Squared Error (RMSE): {rmse_loaded}')
print(f'Mean Absolute Error (MAE): {mae_loaded}')
print(f'R-squared (R2): {r2_loaded}')

```


Mean Squared Error (MSE): 0.2814475768773077
Root Mean Squared Error (RMSE): 0.5305163304529915
Mean Absolute Error (MAE): 0.38660633427178753
R-squared (R²): 0.7171190995832423

Scatter plot of actual vs predicted values

```
In [66]: plt.style.use('dark_background')

# Scatter plot of actual vs predicted values

plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred_loaded, color='blue', s=80, alpha=1, edge
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red'
plt.xlabel('Actual Values', fontsize=14)
plt.ylabel('Predicted Values', fontsize=14)
plt.title('Actual vs Predicted Values', fontsize=16)
plt.legend(fontsize=12)
plt.grid(True)
plt.show()
```



Generating Future Data with Month-Based Historical Averages

```
In [67]: # Extract the last OBSERVATION_DATE from final_merge

last_date = pd.to_datetime(final_merge['OBSERVATION_DATE']).max()

# Generate the next 6 months' dates from the last OBSERVATION_DATE

future_dates = pd.date_range(last_date + timedelta(days=30), periods=6, freq

# Create future_data with OBSERVATION_DATE, Year, and Month columns
```

```

future_data = pd.DataFrame({
    'OBSERVATION_DATE': future_dates,
    'Year': future_dates.year,
    'Month': future_dates.month
})

# Calculate month-based historical means for each feature from numeric_data

monthly_means = numeric_data.groupby('Month').mean()

# Populate future_data with month-based means

for feature in ['EVAP', 'RADIATION', 'Temp_Range', 'RH_TMAX', 'MIN_TEMP', 'MAX_TEMP']:
    future_data[feature] = future_data['Month'].map(monthly_means[feature])

# Display future_data

print(future_data)

```

	OBSERVATION_DATE	Year	Month	EVAP	RADIATION	Temp_Range	RH_TMAX	\
0	2023-02-01	2023	2	1.013965	0.730433	0.138253	-0.443325	
1	2023-03-01	2023	3	0.313775	0.185656	0.023387	-0.119301	
2	2023-04-01	2023	4	-0.327623	-0.233968	-0.044214	0.037131	
3	2023-05-01	2023	5	-0.909433	-0.770360	-0.121155	0.429327	
4	2023-06-01	2023	6	-1.087614	-1.135736	-0.451711	0.686555	
5	2023-07-01	2023	7	-0.913735	-1.036689	-0.543589	0.657553	

	MIN_TEMP	MAX_TEMP
0	0.929178	0.849653
1	0.710950	0.594753
2	-0.044463	-0.065915
3	-0.719516	-0.667405
4	-0.908600	-1.043461
5	-0.895355	-1.094392

Predicting Future TSDM_MEAN Using Trained Model on Generated Data

```

In [68]: # Select only the features required for prediction

X_future = future_data[selected_feature]

# Use the trained model to predict TSDM_MEAN for the next 6 months

future_predictions = loaded_rf_model.predict(X_future)

# Add the predictions to the future data DataFrame

future_data['TSDM_MEAN_Predicted'] = future_predictions

# Display the final DataFrame with future dates and predictions

print(future_data[['OBSERVATION_DATE', 'TSDM_MEAN_Predicted']])

```

```

OBSERVATION_DATE    TSDM_MEAN_Predicted
0      2023-02-01          0.569451
1      2023-03-01          0.816637
2      2023-04-01          0.335159
3      2023-05-01          0.723748
4      2023-06-01          0.832515
5      2023-07-01          1.048704

```

In [69]: `print(future_data.columns)`

```

Index(['OBSERVATION_DATE', 'Year', 'Month', 'EVAP', 'RADIATION', 'Temp_Range',
       'RH_TMAX', 'MIN_TEMP', 'MAX_TEMP', 'TSDM_MEAN_Predicted'],
      dtype='object')

```

In [70]: `future_data.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   OBSERVATION_DATE    6 non-null      datetime64[ns]
 1   Year                6 non-null      int32  
 2   Month               6 non-null      int32  
 3   EVAP                6 non-null      float64
 4   RADIATION           6 non-null      float64
 5   Temp_Range           6 non-null      float64
 6   RH_TMAX              6 non-null      float64
 7   MIN_TEMP             6 non-null      float64
 8   MAX_TEMP             6 non-null      float64
 9   TSDM_MEAN_Predicted  6 non-null      float64
dtypes: datetime64[ns](1), float64(7), int32(2)
memory usage: 564.0 bytes

```

Historical vs Predicted Monthly Mean TSDM_MEAN

In [71]: `# Convert 'OBSERVATION_DATE' to datetime`

```

final_merge['OBSERVATION_DATE'] = pd.to_datetime(final_merge['OBSERVATION_DATE'])
future_data['OBSERVATION_DATE'] = pd.to_datetime(future_data['OBSERVATION_DATE'])

# Set 'OBSERVATION_DATE' as the index for historical data
final_merge.set_index('OBSERVATION_DATE', inplace=True)

# Calculate monthly mean for TSDM_MEAN from historical data
monthly_mean_tsdm = final_merge['TSDM_MEAN'].resample('M').mean()

future_data.set_index('OBSERVATION_DATE', inplace=True)

monthly_predicted_tsdm = future_data['TSDM_MEAN_Predicted'].resample('M').mean()

# Plotting

```

```

plt.figure(figsize=(15, 8), facecolor='black')

# Plotting Historical Monthly Mean TSDM_MEAN

plt.plot(monthly_mean_tsdm.index, monthly_mean_tsdm,
         color='blue', label='Historical Monthly Mean TSDM_MEAN', linewidth=2)

# Plotting Predicted Monthly Mean TSDM_MEAN


plt.plot(monthly_predicted_tsdm.index, monthly_predicted_tsdm,
         color='orange', linestyle='--', label='Predicted Monthly Mean TSDM_MEAN', linewidth=2)

plt.title("Historical vs Predicted Monthly Mean TSDM_MEAN", color='white', fontweight='bold')
plt.xlabel("Date", color='white', fontsize=14)
plt.ylabel("TSDM_MEAN", color='white', fontsize=14)
plt.legend(frameon=True, loc='upper left', fontsize='small', facecolor='black')
plt.grid(color='grey', linestyle='--', linewidth=0.5)
plt.tight_layout()
plt.show()

# Display the monthly means DataFrames

print("Monthly Mean TSDM_MEAN:")
print(monthly_mean_tsdm)
print("\nMonthly Predicted TSDM_MEAN:")
print(monthly_predicted_tsdm)

```



Monthly Mean TSDM_MEAN:

OBSERVATION_DATE	TSDM_MEAN
2017-01-31	0.263005
2017-02-28	0.255832
2017-03-31	0.193568
2017-04-30	0.215812
2017-05-31	0.195135
	...
2022-08-31	0.803227
2022-09-30	0.783521
2022-10-31	0.921857
2022-11-30	1.134149
2022-12-31	1.018216

Freq: M, Name: TSDM_MEAN, Length: 72, dtype: float64

Monthly Predicted TSDM_MEAN:

OBSERVATION_DATE	TSDM_MEAN_Predicted
2023-02-28	0.569451
2023-03-31	0.816637
2023-04-30	0.335159
2023-05-31	0.723748
2023-06-30	0.832515
2023-07-31	1.048704

Freq: M, Name: TSDM_MEAN_Predicted, dtype: float64